

## Report

### Approach-

**Data Preprocessing:** The dataset information is read from a CSV file (train.csv) containing image filenames and class labels. The data is then split into training and validation sets using `train_test_split` from `sklearn`. For image augmentation, the `ImageDataGenerator` is used: the training data is augmented with rescaling, rotation, width/height shift, shear, zoom, and horizontal flip to enhance variety and prevent overfitting, while the validation data is only rescaled. Image data generators are created for both training and validation sets using `flow_from_dataframe`, which reads filenames and class labels from the `DataFrame` and fetches the images from specified directories.

**Model Definition:** The base model used is `MobileNetV2` with pre-trained `ImageNet` weights, excluding the top classification layers (`include_top=False`) to allow customization. The output of the base model is flattened, followed by a fully connected (`Dense`) layer with 512 neurons and `ReLU` activation. A `Dropout` layer with a 0.5 rate is added to prevent overfitting. The final layer is a `Dense` layer with neurons equal to the number of unique classes, using `softmax` activation for multi-class classification. All base model layers are frozen to keep their weights unchanged during training. The model is compiled with the `Adam` optimizer (learning rate  $1e-4$ ) and `categorical_crossentropy` as the loss function, with accuracy as the evaluation metric.

**Model Training:** The model is trained for 80 epochs using the `fit` method with training and validation data generators. The training history, which includes metrics like training and validation accuracy and loss, is stored for further analysis.

**Testing and Predictions:** Test image filenames are collected from the test directory and stored in a `DataFrame`. An `ImageDataGenerator` is created for the test data, applying only rescaling. The model predicts class probabilities for the test images, and predicted class labels are obtained using `argmax`. These predictions are saved to a CSV file (`KBLD_sbmission.csv`).

### Accuracy and Loss-

Accuracy = 88% Loss = 0.36 Epochs = 5 Accuracy = 95.21% Loss = 0.13 Epochs = 80